

Clustering: Mixture Models

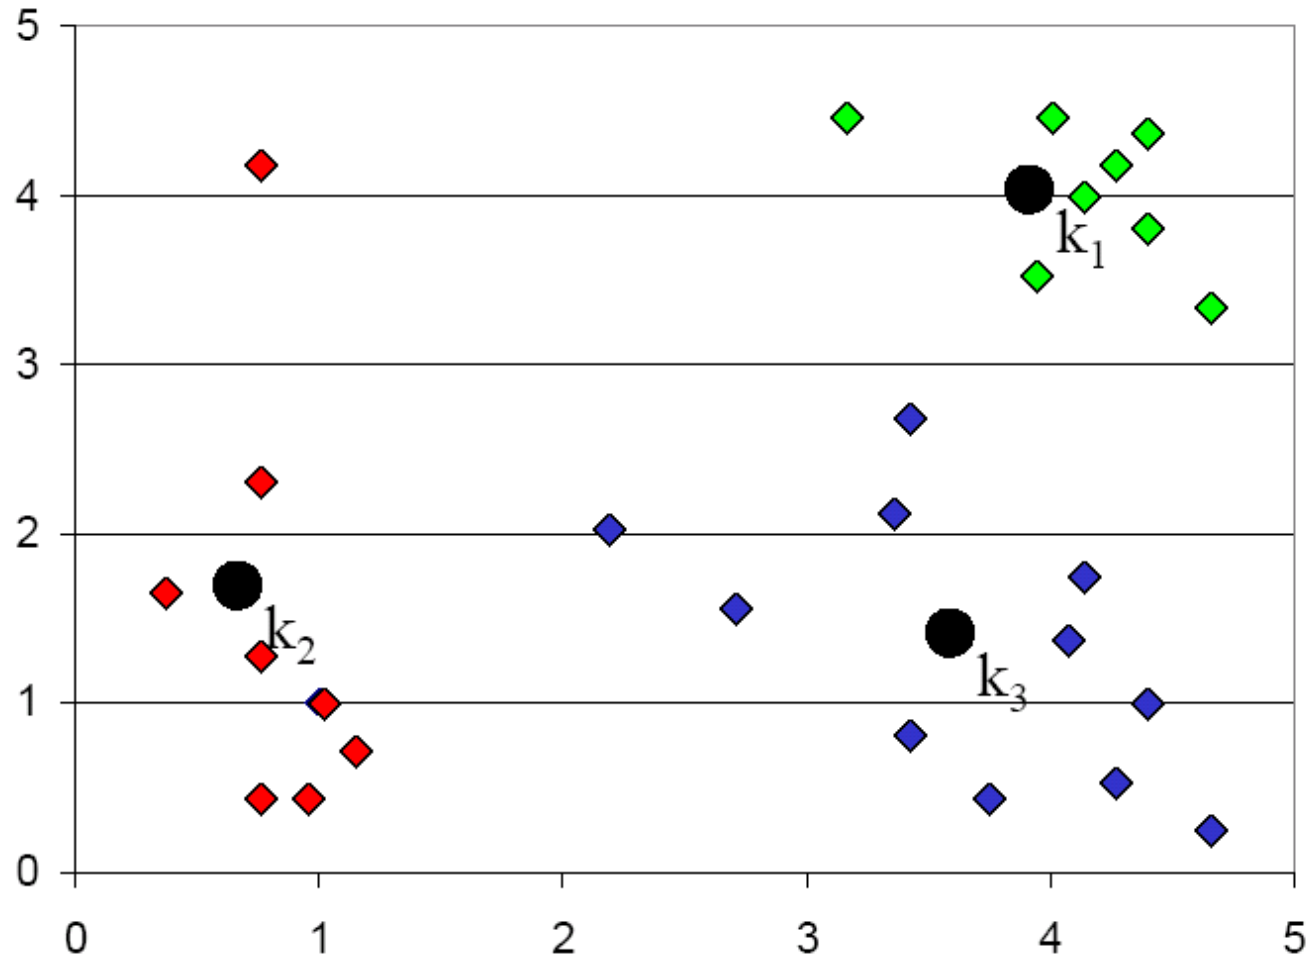
Machine Learning 10-601B

Seyoung Kim

Many of these slides are derived from Tom Mitchell, Ziv-Bar Joseph, and Eric Xing. Thanks!



Problem with K-means



Hard Assignment of Samples into Three Clusters

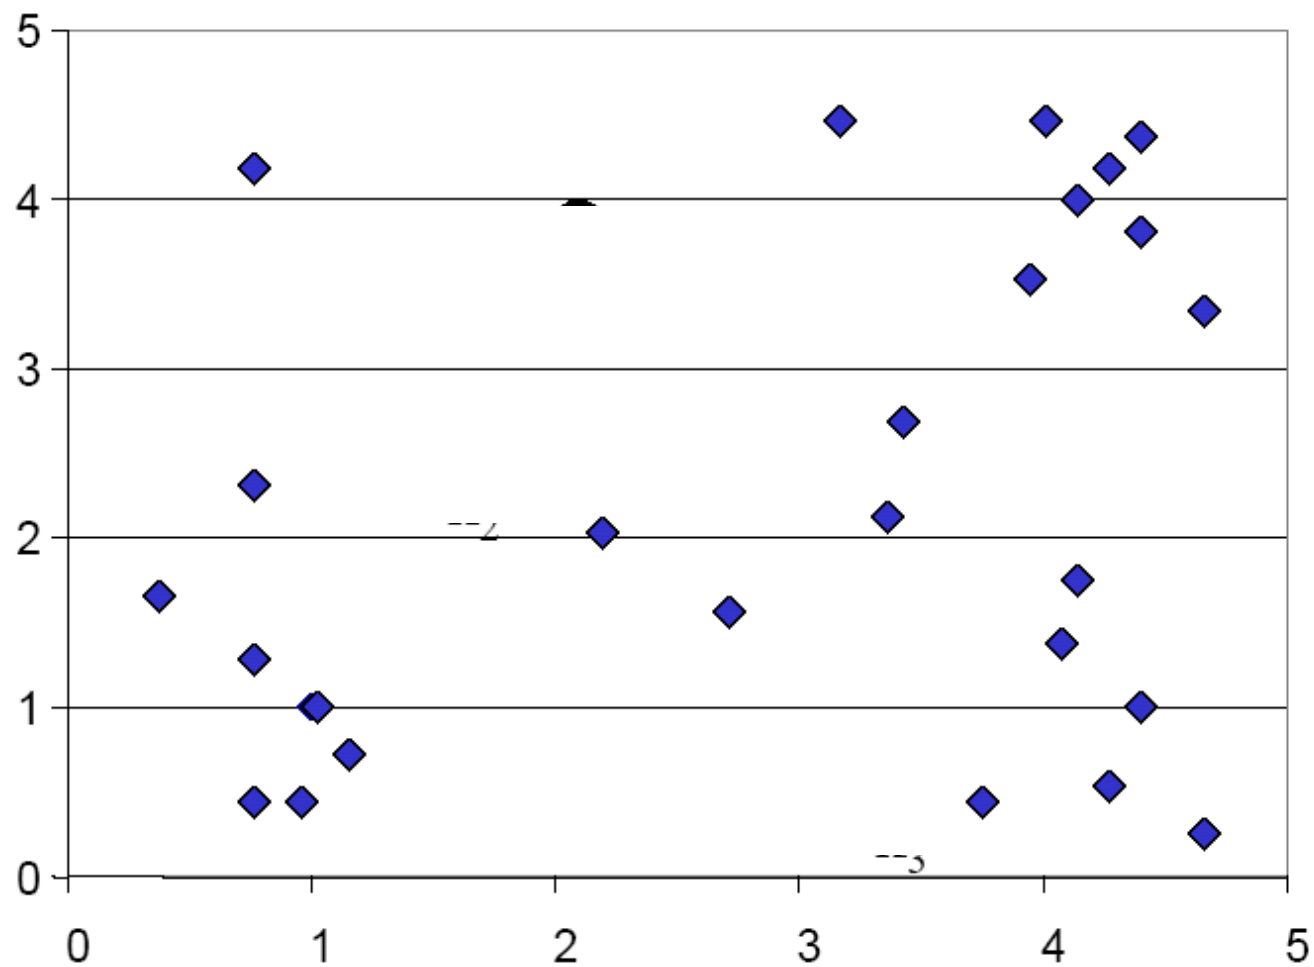
	Cluster 1	Cluster 2	Cluster 3
Individual 1	1	0	0
Individual 2	0	1	0
Individual 3	0	1	0
Individual 4	1	0	0
Individual 5
Individual 6
Individual 7
Individual 8
Individual 9
Individual 10

Probabilistic Soft-Clustering of Samples into Three Clusters

Probability of	Cluster 1	Cluster 2	Cluster 3	Sum
Individual 1	0.1	0.4	0.5	1
Individual 2	0.8	0.1	0.1	1
Individual 3	0.7	0.2	0.1	1
Individual 4	0.10	0.05	0.85	1
Individual 5	1
Individual 6	1
Individual 7	1
Individual 8	1
Individual 9	1
Individual 10	1

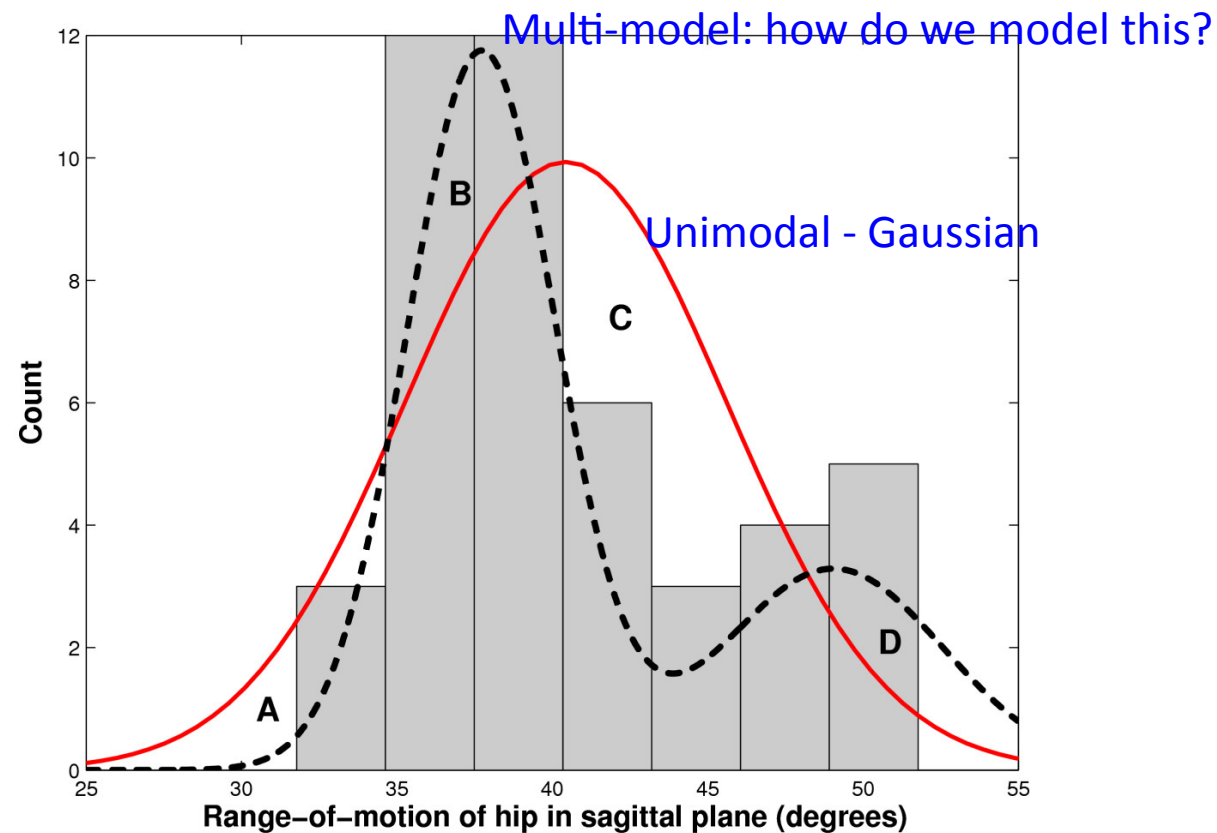
- Each sample can be assigned to more than one clusters with a certain probability.
- For each sample, the probabilities for all clusters should sum to 1. (i.e., each row should sum to 1.)
- Each cluster is explained by a cluster center variable (i.e., cluster mean)

Probability Model for Data $P(X)$?



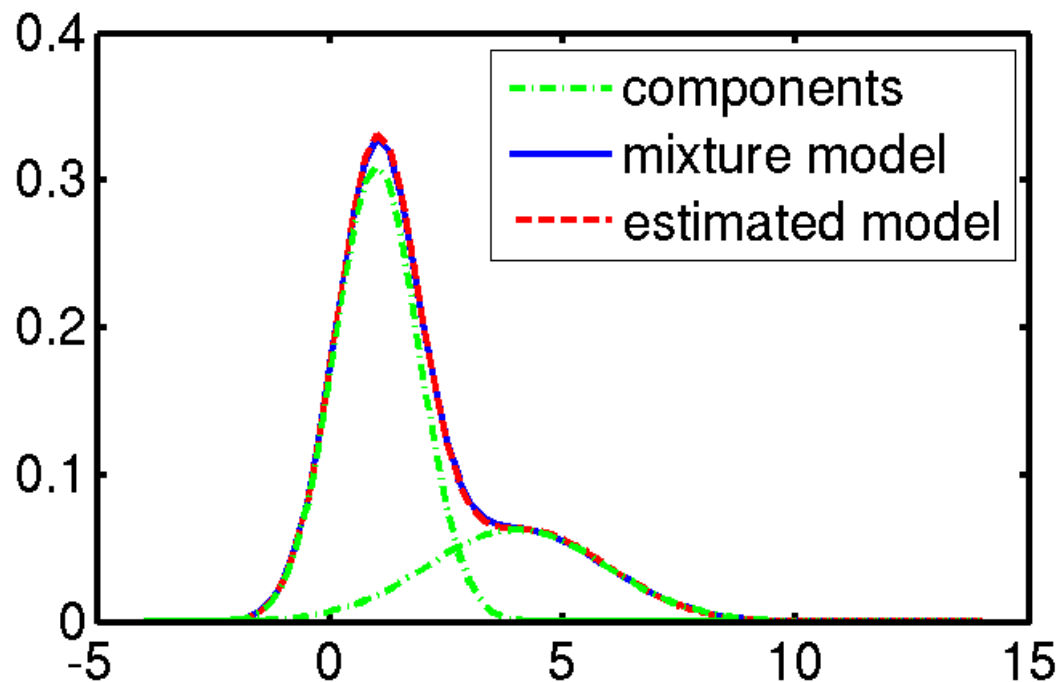
Mixture Model

- A density model $p(\mathbf{x})$ may be multi-modal.



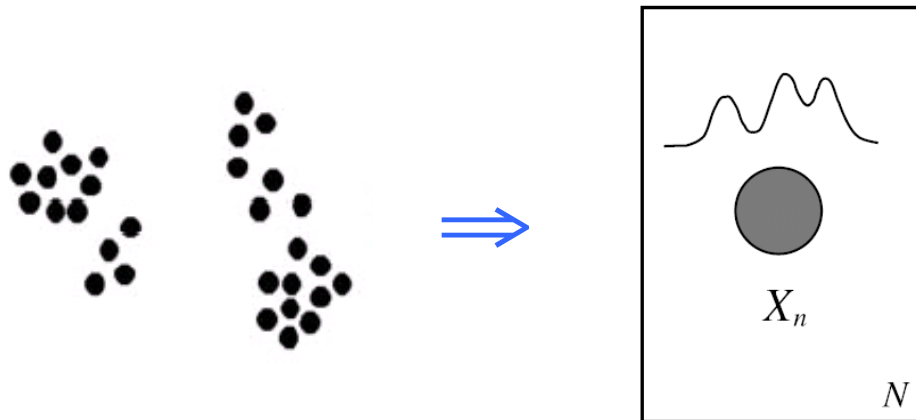
Mixture Model

- We may be able to model it as a mixture of uni-modal distributions (e.g., Gaussians).
- Each mode may correspond to a different sub-population (e.g., male and female).

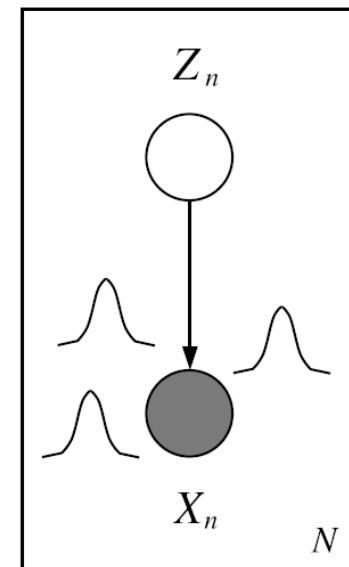


Learning Mixture Models from Data

- Given data generated from multi-modal distribution, can we find a representation of the multi-modal distribution as a mixture of uni-modal distributions?



(a)



(b)

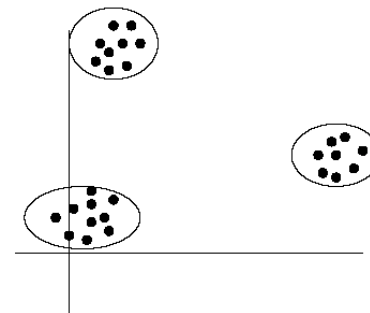
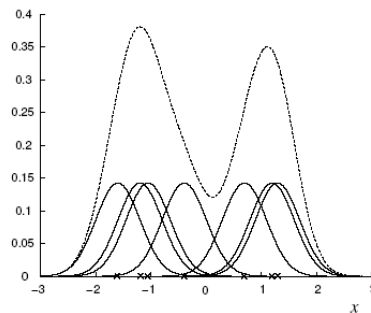
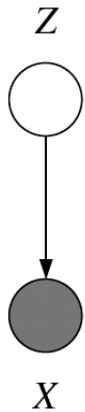
Gaussian Mixture Models (GMMs)

- Consider a mixture of K Gaussian components:

$$p(x_n) = \sum_k p(x_n | z_n = k) p(z_n = k)$$

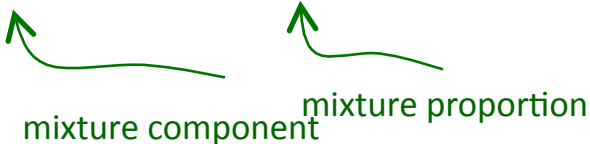
$$= \sum_k N(x_n | \mu_k, \Sigma_k) \pi_k$$

mixture component mixture proportion



Gaussian Mixture Models (GMMs)

- Consider a mixture of K Gaussian components:

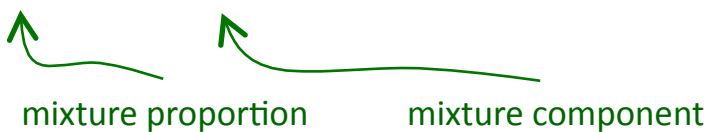
$$\begin{aligned} p(x_n) &= \sum_k p(x_n, z_n = k) p(z_n = k) \\ &= \sum_k N(x_n | \mu_k, \Sigma_k) \pi_k \end{aligned}$$


mixture component mixture proportion

- This probability model describes how each data point x_n can be generated
 - Step 1: Flip a K -sided die (with probability π_k for the k -th side) to select a cluster c
 - Step 2: Generate the values of the data point from $N(\mu_c, \Sigma_c)$

Gaussian Mixture Models (GMMs)

- Consider a mixture of K Gaussian components:

$$\begin{aligned} p(x_n) &= \sum_k p(x_n, z_n = k) p(z_n = k) \\ &= \sum_k N(x_n | \mu_k, \Sigma_k) \pi_k \end{aligned}$$


mixture proportion mixture component

- Parameters for K clusters: $\theta = \{\mu_k, \Sigma_k, \pi_k, k = 1, \dots, K\}$

Learning mixture models

- Latent variable model: data are only partially observed!
 - x_i : **observed** sample data
 - $z_i = \{z_i^1 \dots z_i^K\}$: **Unobserved** cluster labels (each element 0 or 1, only one of them is 1)
- MLE estimate
 - What if all data (x_i, z_i) are observed?
 - Maximize the data log likelihood for (x_i, z_i) based on $p(x_i, z_i)$
 - **Easy** to optimize!
 - In practice, only x_i 's are observed
 - Maximize the data log likelihood for (x_i) based on $p(x_i)$
 - **Difficult** to optimize!
 - Maximize the **expected** data log likelihood for (x_i, z_i) based on $p(x_i, z_i)$
 - Expectation-Maximization (EM) algorithm

Learning mixture models: fully observed data

- In fully observed iid settings, assuming the **cluster labels z_i 's were observed**, the log likelihood decomposes into a sum of local terms.

$$l_c(\theta; D) = \sum_n \log p(x_n, z_n | \theta) = \underbrace{\sum_n \log p(z_n | \theta)}_{\text{Depends on } \pi_k} + \underbrace{\sum_n \log p(x_n | z_n, \theta)}_{\text{Depends on } \mu_k, \Sigma_k}$$

- The optimization problems for μ_k, Σ_k and for π_k are decoupled, and a closed-form solution for MLE exists.

MLE for GMM with fully observed data

- If we are doing MLE for **completely observed data**

- Data log-likelihood

$$\begin{aligned}
 l(\theta; D) &= \log \prod_n p(z_n, x_n) = \log \prod_n p(z_n | \pi) p(x_n | z_n, \mu, \sigma) \\
 &= \sum_n \log \prod_k \pi_k^{z_n^k} + \sum_n \log \prod_k N(x_n; \mu_k, \sigma)^{z_n^k} \\
 &= \sum_n \sum_k z_n^k \log \pi_k - \sum_n \sum_k z_n^k \frac{1}{2\sigma^2} (x_n - \mu_k)^2 + C
 \end{aligned}$$

- MLE

$$\hat{\pi}_{k,MLE} = \arg \max_{\pi} l(\theta; D),$$

$$\hat{\mu}_{k,MLE} = \arg \max_{\mu} l(\theta; D)$$

$$\hat{\sigma}_{k,MLE} = \arg \max_{\sigma} l(\theta; D)$$

$$\hat{\mu}_{k,MLE} = \frac{\sum_n z_n^k x_n}{\sum_n z_n^k}$$

- What if we do not know z_n ?

Learning mixture models

- In fully observed iid settings, assuming the cluster labels z_i 's were observed, the log likelihood decomposes into a sum of local terms.

$$l_c(\theta; D) = \sum_n \log p(x_n, z_n | \theta)$$

- With latent variables for cluster labels

$$\begin{aligned} l_c(\theta; D) &= \sum_n \log p(x_n | \theta) \\ &= \sum_n \log \sum_z p(x_n, z | \theta) = \sum_n \log \sum_z p(z | \theta) p(x_n | z, \theta) \end{aligned}$$

– all the parameters become coupled together via *marginalization*

- Are they equally difficult?

Depends on π_k

Depends on μ_k, Σ_k

Theory underlying EM

- Recall that according to MLE, we intend to learn the model parameter that would have maximized the likelihood of the data.
- But we do not observe z , so computing

$$l_c(\theta; D) = \sum_n \log \sum_z p(x_n, z | \theta) = \sum_n \log \sum_z p(z | \theta) p(x_n | z, \theta)$$

is difficult!

- Optimizing the log-likelihood for MLE is difficult!
- What shall we do?

Complete vs. Expected Complete Log Likelihoods

- The complete log likelihood:

$$\begin{aligned}
 l(\theta; D) &= \log \prod_n p(z_n, x_n) = \log \prod_n p(z_n | \pi) p(x_n | z_n, \mu, \sigma) \\
 &= \sum_n \log \prod_k \pi_k^{z_n^k} + \sum_n \log \prod_k N(x_n; \mu_k, \sigma)^{z_n^k} \\
 &= \sum_n \sum_k z_n^k \log \pi_k - \sum_n \sum_k z_n^k \frac{1}{2\sigma^2} (x_n - \mu_k)^2 + C
 \end{aligned}$$

- The expected complete log likelihood

$$\begin{aligned}
 \langle l_c(\theta; \mathbf{x}, \mathbf{z}) \rangle &= \sum_n \langle \log p(\mathbf{z}_n | \pi) \rangle_{p(\mathbf{z}|\mathbf{x})} + \sum_n \langle \log p(\mathbf{x}_n | \mathbf{z}_n, \mu, \Sigma) \rangle_{p(\mathbf{z}|\mathbf{x})} \\
 &= \sum_n \sum_k \langle z_n^k \rangle \log \pi_k - \frac{1}{2} \sum_n \sum_k \langle z_n^k \rangle \left((\mathbf{x}_n - \mu_k)^T \Sigma_k^{-1} (\mathbf{x}_n - \mu_k) + \log |\Sigma_k| + C \right)
 \end{aligned}$$

Depends on π_k
Depends on μ_k, Σ_k

Complete vs. Expected Complete Log Likelihoods

- The complete log likelihood:

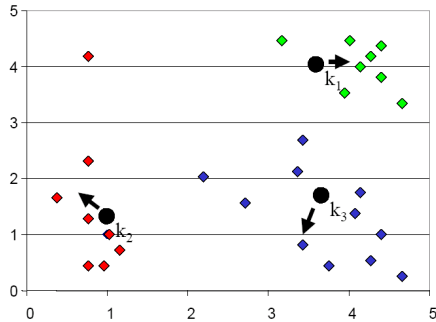
$$\begin{aligned}
 l(\theta; D) &= \log \prod_n p(z_n, x_n) = \log \prod_n p(z_n | \pi) p(x_n | z_n, \mu, \sigma) \\
 &= \sum_n \log \prod_k \pi_k^{z_n^k} + \sum_n \log \prod_k N(x_n; \mu_k, \sigma)^{z_n^k} \\
 &= \sum_n \sum_k z_n^k \log \pi_k - \sum_n \sum_k z_n^k \frac{1}{2\sigma^2} (x_n - \mu_k)^2 + C
 \end{aligned}$$

- The expected complete log likelihood

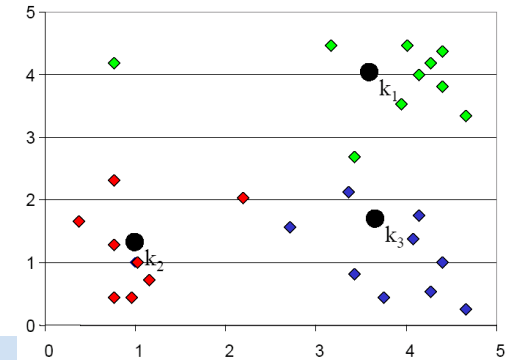
$$\begin{aligned}
 \langle l_c(\theta; \mathbf{x}, \mathbf{z}) \rangle &= \sum_n \langle \log p(\mathbf{z}_n | \pi) \rangle_{p(\mathbf{z}|\mathbf{x})} + \sum_n \langle \log p(\mathbf{x}_n | \mathbf{z}_n, \mu, \Sigma) \rangle_{p(\mathbf{z}|\mathbf{x})} \\
 &= \sum_n \sum_k \langle z_n^k \rangle \log \pi_k - \frac{1}{2} \sum_n \sum_k \langle z_n^k \rangle \left((\mathbf{x}_n - \mu_k)^T \Sigma_k^{-1} (\mathbf{x}_n - \mu_k) + \log |\Sigma_k| + C \right)
 \end{aligned}$$

- EM optimizes the expected complete log likelihood

EM Algorithm



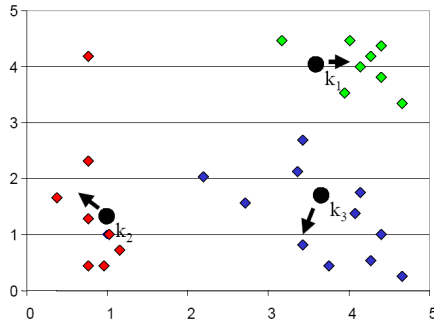
Maximization (M)-step:
- Find mixture parameters



Expectation (E)-step:
- Re-assign samples x_i 's to clusters
- Impute the unobserved values z_i

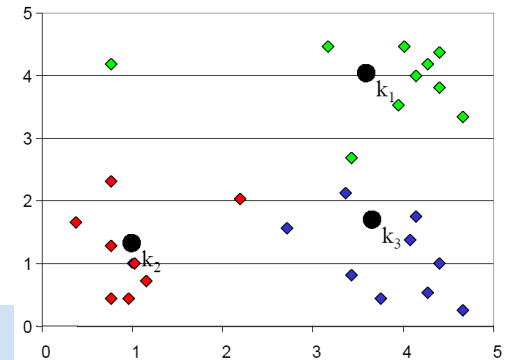
Iterate until convergence

K-Means Clustering Algorithm



Find the cluster means

$$\vec{\mu}_k = \frac{1}{c_k} \sum_{i \in \mathcal{C}_k} \vec{x}_i$$



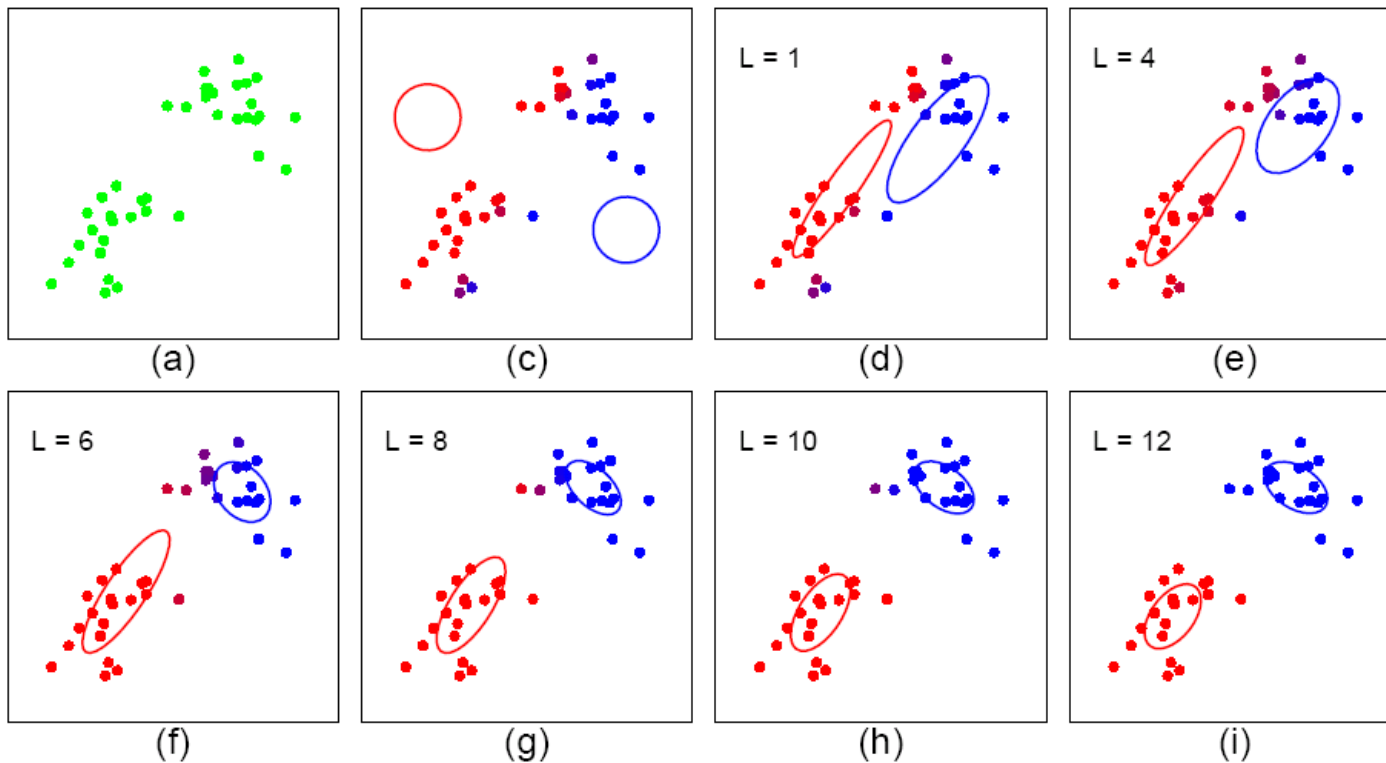
Re-assign samples x_i 's to clusters

$$\operatorname{argmax}_k \|x_i - \mu_k\|_2^2$$

Iterate until convergence

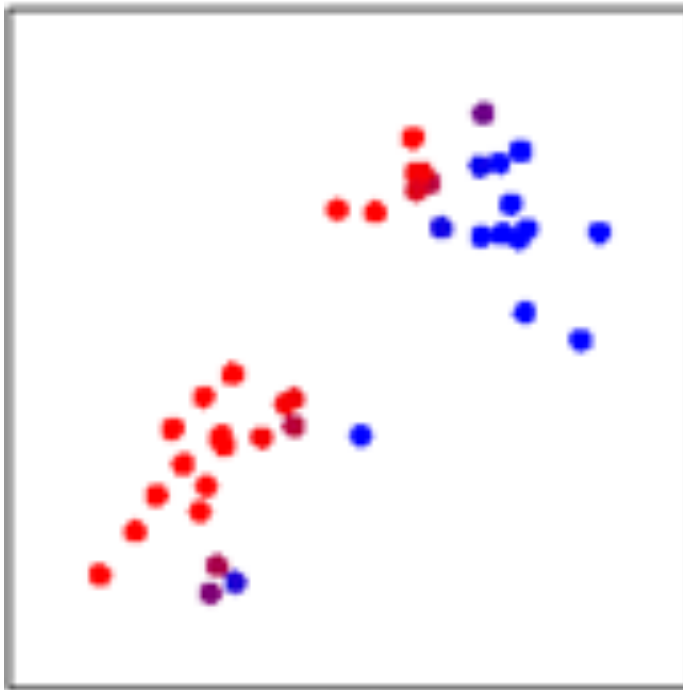
The Expectation-Maximization (EM) Algorithm

- Start:
 - "Guess" the centroid μ_k and covariance Σ_k of each of the K clusters
- Loop



The Expectation-Maximization (EM) Algorithm

- A “soft” k-means



E step:

$$\tau_n^{k(t)} = \langle z_n^k \rangle_{q^{(t)}} = p(z_n^k = 1 \mid x, \mu^{(t)}, \Sigma^{(t)})$$

M step:

$$\pi_k^{(t+1)} = \frac{\sum_n \tau_n^{k(t)}}{N} = \frac{\langle n_k \rangle}{N}$$

$$\mu_k^{(t+1)} = \frac{\sum_n \tau_n^{k(t)} x_n}{\sum_n \tau_n^{k(t)}}$$

$$\Sigma_k^{(t+1)} = \frac{\sum_n \tau_n^{k(t)} (x_n - \mu_k^{(t+1)})(x_n - \mu_k^{(t+1)})^T}{\sum_n \tau_n^{k(t)}}$$

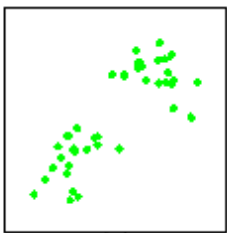
Compare: K-means

- The EM algorithm for mixtures of Gaussians is like a "soft version" of the K-means algorithm.
- In the K-means “E-step” we do hard assignment:

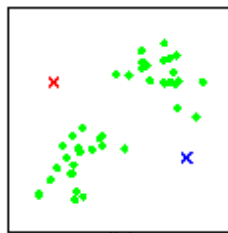
$$\mathbf{z}_n^{(t)} = \arg \max_k (\mathbf{x}_n - \mu_k^{(t)})^T \Sigma_k^{-1(t)} (\mathbf{x}_n - \mu_k^{(t)})$$

- In the K-means “M-step” we update the means as the weighted sum of the data, but now the weights are 0 or 1:

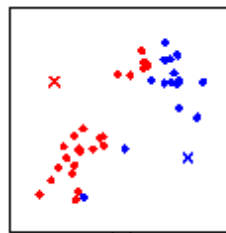
$$\mu_k^{(t+1)} = \frac{\sum_n \delta(\mathbf{z}_n^{(t)}, k) \mathbf{x}_n}{\sum_n \delta(\mathbf{z}_n^{(t)}, k)}$$



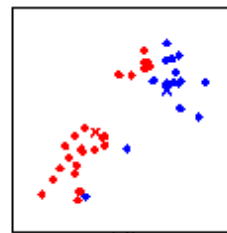
(a)



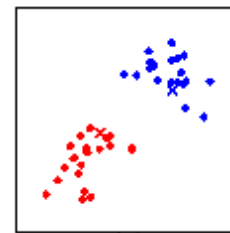
(b)



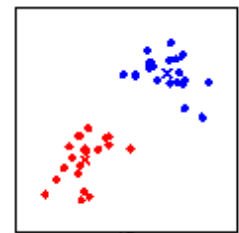
(c)



(d)



(e)



(f)

Expected Complete Log Likelihood Lower-bounds

Complete Log Likelihood

- For **any** distribution $q(\mathbf{z})$, define *expected complete log likelihood*:

$$\langle I_c(\theta; \mathbf{x}, \mathbf{z}) \rangle_q \stackrel{\text{def}}{=} \sum_{\mathbf{z}} q(\mathbf{z} | \mathbf{x}, \theta) \log p(\mathbf{x}, \mathbf{z} | \theta)$$

– Does maximizing this surrogate yield a maximizer of the likelihood?

- Jensen's inequality

$$\begin{aligned} I(\theta; \mathbf{x}) &= \log p(\mathbf{x} | \theta) \\ &= \log \sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z} | \theta) \\ &= \log \sum_{\mathbf{z}} q(\mathbf{z} | \mathbf{x}) \frac{p(\mathbf{x}, \mathbf{z} | \theta)}{q(\mathbf{z} | \mathbf{x})} \\ &\geq \sum_{\mathbf{z}} q(\mathbf{z} | \mathbf{x}) \log \frac{p(\mathbf{x}, \mathbf{z} | \theta)}{q(\mathbf{z} | \mathbf{x})} \quad \Rightarrow \quad I(\theta; \mathbf{x}) \geq \langle I_c(\theta; \mathbf{x}, \mathbf{z}) \rangle_q + H_q \end{aligned}$$

Closing notes

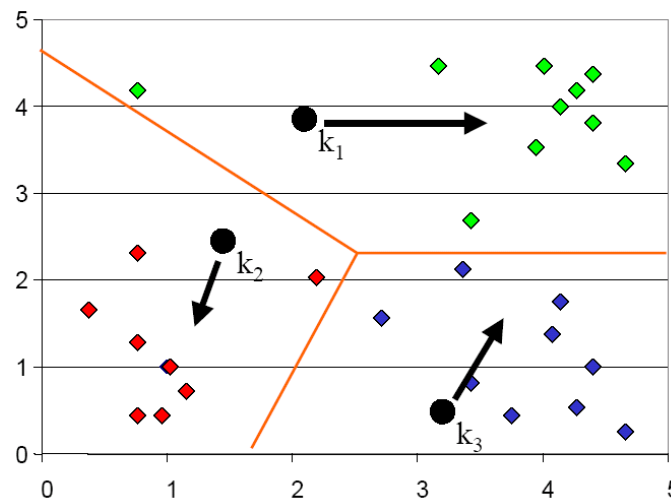
- Convergence
- Seed choice
- Quality of cluster
- How many clusters

Convergence

- Why should the K-means algorithm ever reach a fixed point?
 - -- A state in which clusters don't change.
- K-means is a special case of a general procedure the Expectation Maximization (EM) algorithm.
 - Both are known to converge.
 - Number of iterations could be large.

Seed Choice

- Results can vary based on random seed selection.



- Some seeds can result in convergence to sub-optimal clusterings.
 - Select good seeds using a heuristic (e.g., doc least similar to any existing mean)
 - Try out multiple starting points (very important!!!)
 - Initialize with the results of another method.

What Is A Good Clustering?

- Internal criterion: A good clustering will produce high quality clusters in which:
 - the intra-class (that is, intra-cluster) similarity is high
 - the inter-class similarity is low
 - The measured quality of a clustering depends on both the obj representation and the similarity measure used
- External criteria for clustering quality
 - Quality measured by its ability to discover some or all of the hidden patterns or latent classes in gold standard data
 - Assesses a clustering with respect to ground truth

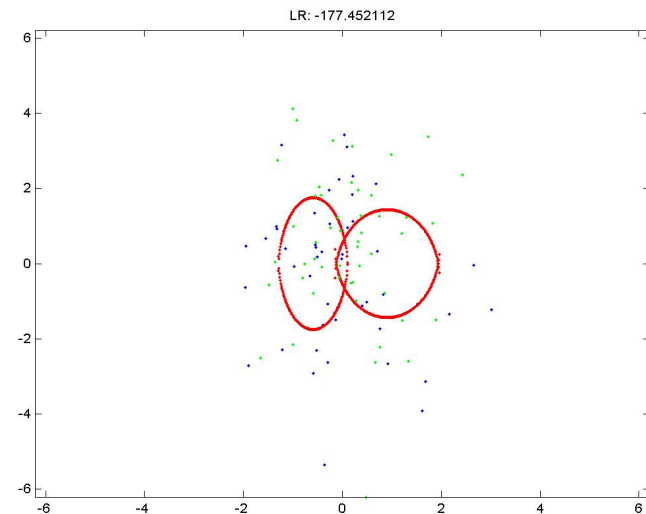
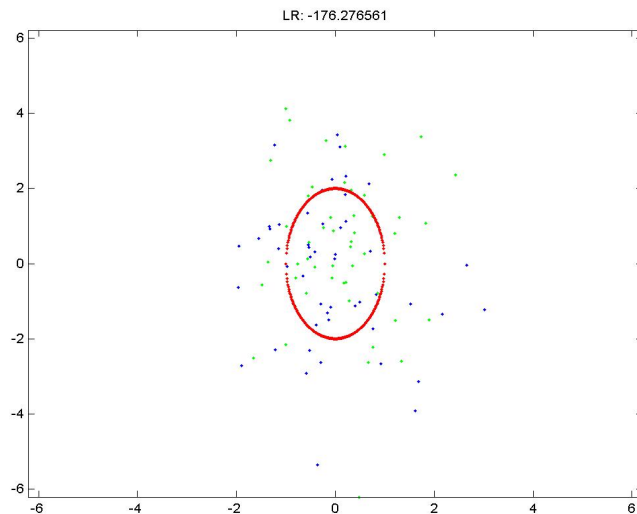
How Many Clusters?

- Number of clusters K is given
 - Partition n docs into predetermined number of clusters
- Finding the “right” number of clusters is part of the problem
 - Given objs, partition into an “appropriate” number of subsets.
 - E.g., for query results - ideal value of K not known up front - though UI may impose limits.
- Tradeoff between having more clusters (better focus within each cluster) and having too many clusters
- Nonparametric Bayesian Inference

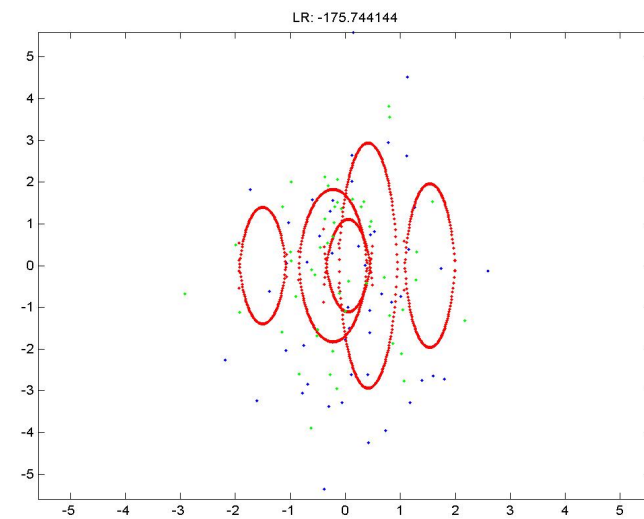
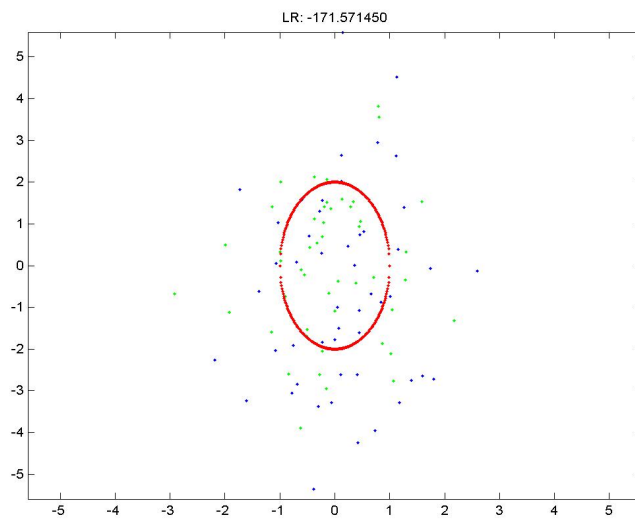
Cross validation

- We can also use cross validation to determine the correct number of classes
- Recall that GMMs is a generative model. We can compute the likelihood of the held-out data to determine which model (number of clusters) is more accurate

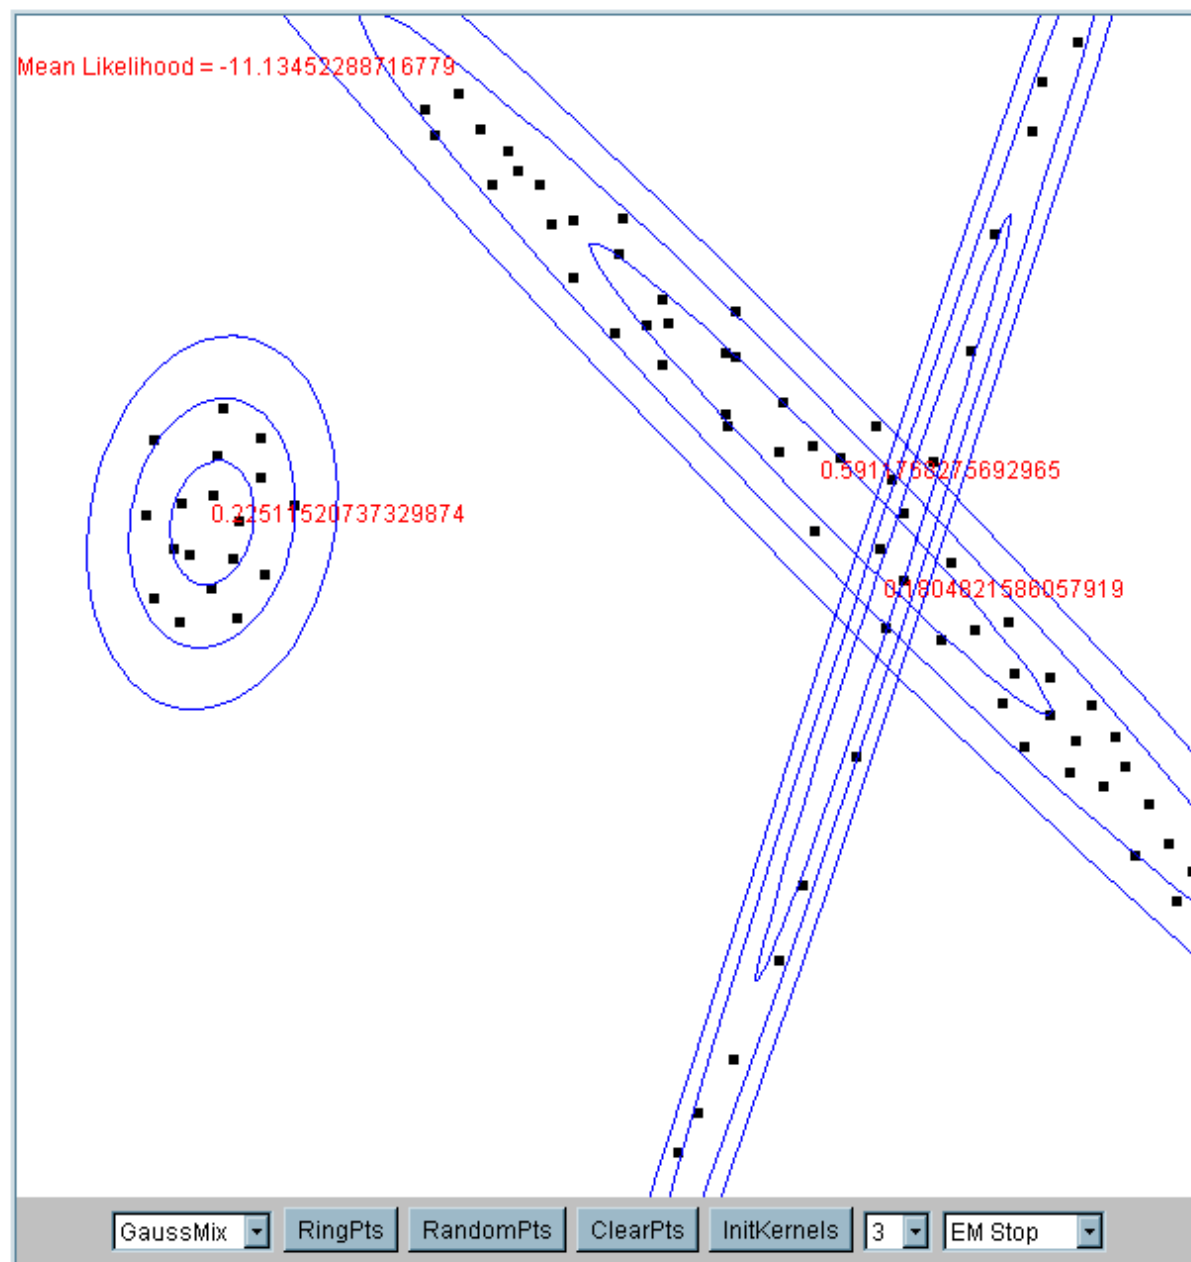
$$p(x_1 \cdots x_n \mid \theta) = \prod_{j=1}^n \left(\sum_{i=1}^k p(x_j \mid C = i) w_i \right)$$



Cross validation



Gaussian mixture clustering



Clustering methods: Comparison

	Hierarchical	K-means	GMM
Running time	naively, $O(N^3)$	fastest (each iteration is linear)	fast (each iteration is linear)
Assumptions	requires a similarity / distance measure	strong assumptions	strongest assumptions
Input parameters	none	K (number of clusters)	K (number of clusters)
Clusters	subjective (only a tree is returned)	exactly K clusters	exactly K clusters

What you should know about Mixture Models

- Gaussian mixture models
 - Probabilistic extension of K-means for soft-clustering
 - EM algorithm for learning by assuming data are only partially observed
 - Cluster labels are treated as the unobserved part of data
- EM algorithm for learning from partly unobserved data
 - MLE of $\theta = \arg \max_{\theta} \log P(data|\theta)$
 - EM estimate: $\theta = \arg \max_{\theta} E_{Z|X,\theta}[\log P(X, Z|\theta)]$
 - Where X is observed part of data, Z is unobserved